

Operating Systems: Introduction

Neerja Mhaskar

Department of Computing and Software, McMaster University, Canada

Acknowledgements: Material based on the textbook Operating Systems Concepts (Chapter 1)

Course Introduction

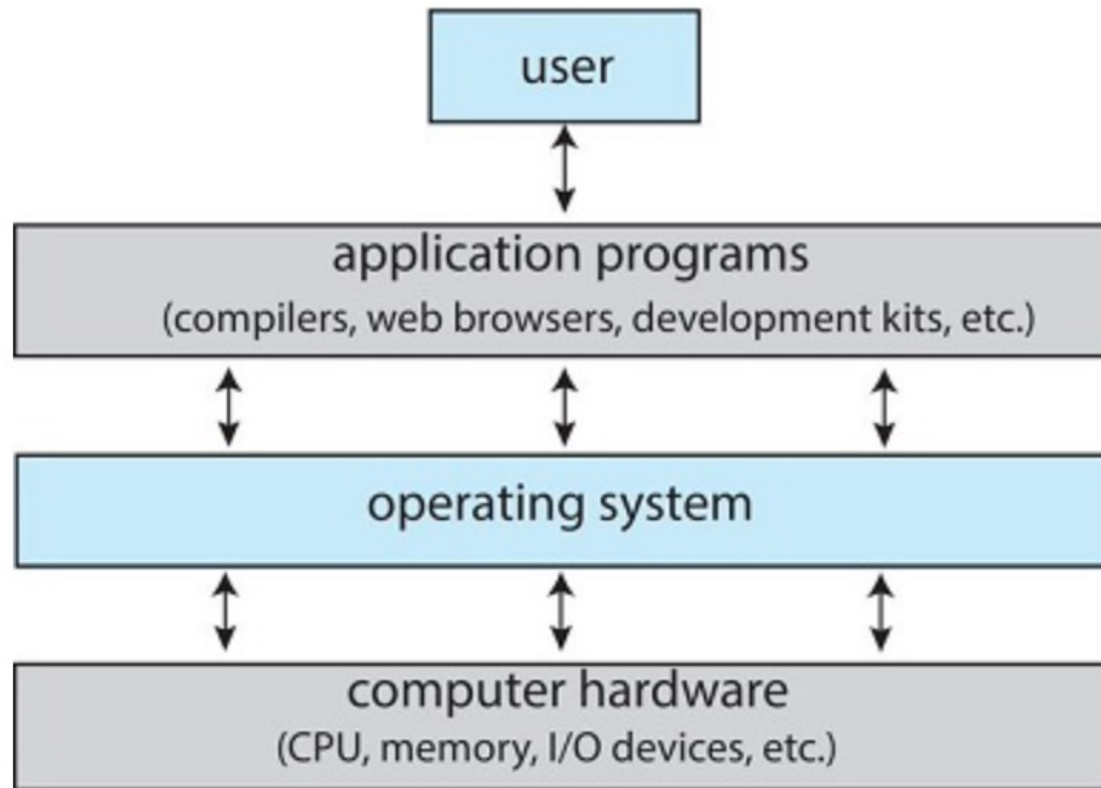
- Take home Lab
 - Installing Virtual Machine
 - Dev Containers
- Groups
 - You will not be able to view or submit assignments if you don't enroll in a group!

What does an Operating System do?

- An Operating system is a
 - **Resource allocator:** Manages a computer's resources
 - **Control program:** Controls execution of programs.
 - Acts as an intermediary between a user of a computer and the computer hardware.
- Operating system goals are
 - *Efficient use*
 - *User convenience*
 - *Non-interference*

Computer System Structure

- Computer system can be divided into **four component**

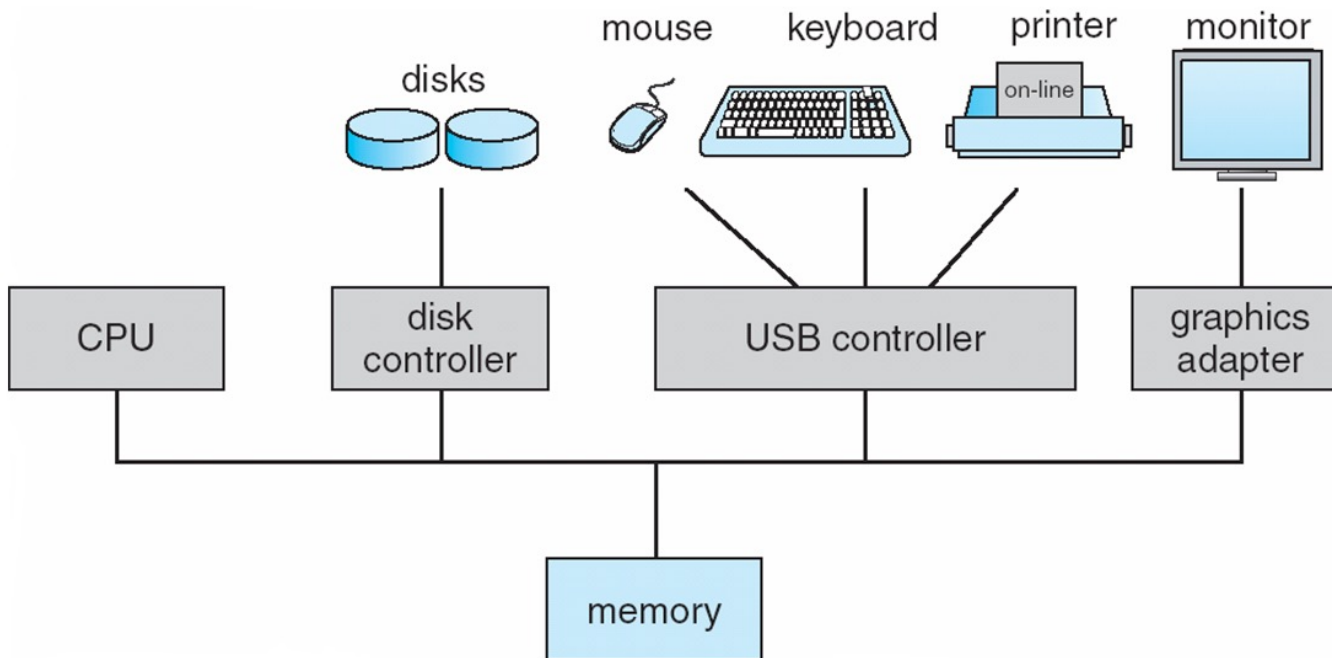


What is an OS?

- Primarily an OS consists of the **Kernel**
 - Kernel stays in main memory (and is the one program that is always running on the computer)
 - Supports process and memory management, some also support file management.
 - Kernel controls the execution of all other programs
 - Other programs (system or user) interact with the kernel through *system calls* (are routines mostly written in a high-level language (C or C++). However, lower-level task written in assembly.)

Computer System Organization

- One or more CPUs, device controllers connect through common bus providing access to shared memory
- Concurrent execution of CPUs and devices competing for memory.



Computer Startup

- When a computer system is switched on or rebooted,
 - it automatically loads a program (**bootstrap program**) stored on a reserved part of an I/O device typically a disk (ROM or EEPROM, generally known as **firmware**) and starts executing the program.
- The bootstrap program then
 - Loads the kernel
 - and system programs (also known as system daemons)
- The kernel and system programs run all the time on the computer to provide services.
- After the system is fully booted it waits for an **event** to occur.

Computer System Operation - Interrupts

- An operating system is **event driven** and events occur by interrupts. Therefore, OS is **interrupt driven**.
- **Interrupt:** is a mechanism that enables a device/software to notify the CPU that it needs attention.
- An interrupt is caused by a
 - **signal to CPU from a device** attached to a computer via system bus (hardware), or
 - from an executing program within the computer through system calls.
- A **trap** or **exception** is a software-generated interrupt caused either by an error or a user request.

Computer System Operation: Interrupt Handling

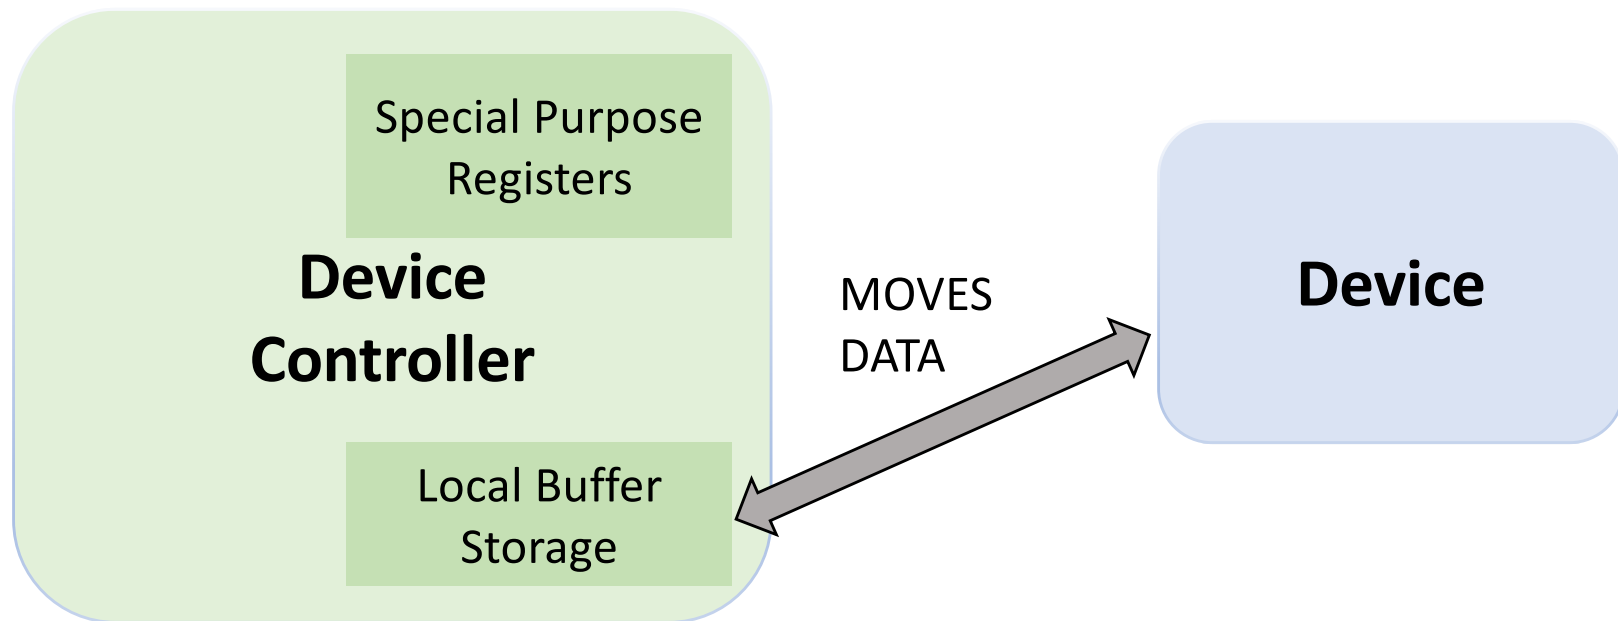
- Interrupt requires the operating system to **stop and figure out what to do next**.
- When an interrupt occurs:
 - CPU stops executing current task
 - The operating system preserves the state of the CPU by storing registers and the program counter
 - Transfers execution to a fixed memory location, which has the starting address of **Interrupt Service Routine (ISR)**
 - ISR handles the interrupt, after which the interrupted process resumes its execution.
- Note: Separate segments of code determine what action should be taken for each type of interrupt

Computer System Operation: Interrupt Handling Implementation

- Implementing ISR as a routine is slow and therefore inefficient.
- Only predefined interrupts exist
 - Use a table of pointers to the various interrupt routines instead!
 - This table/array of addresses is called **Interrupt vector**
- Interrupt vector is usually stored in **low memory**.
- Windows and Linux dispatch interrupts in this manner.


Computer System Operation: I/O Structure

- A general-purpose computer system consists of CPUs and multiple device controllers that are connected through a common bus.
- **Device Controller:** The I/O managing processor within a device.



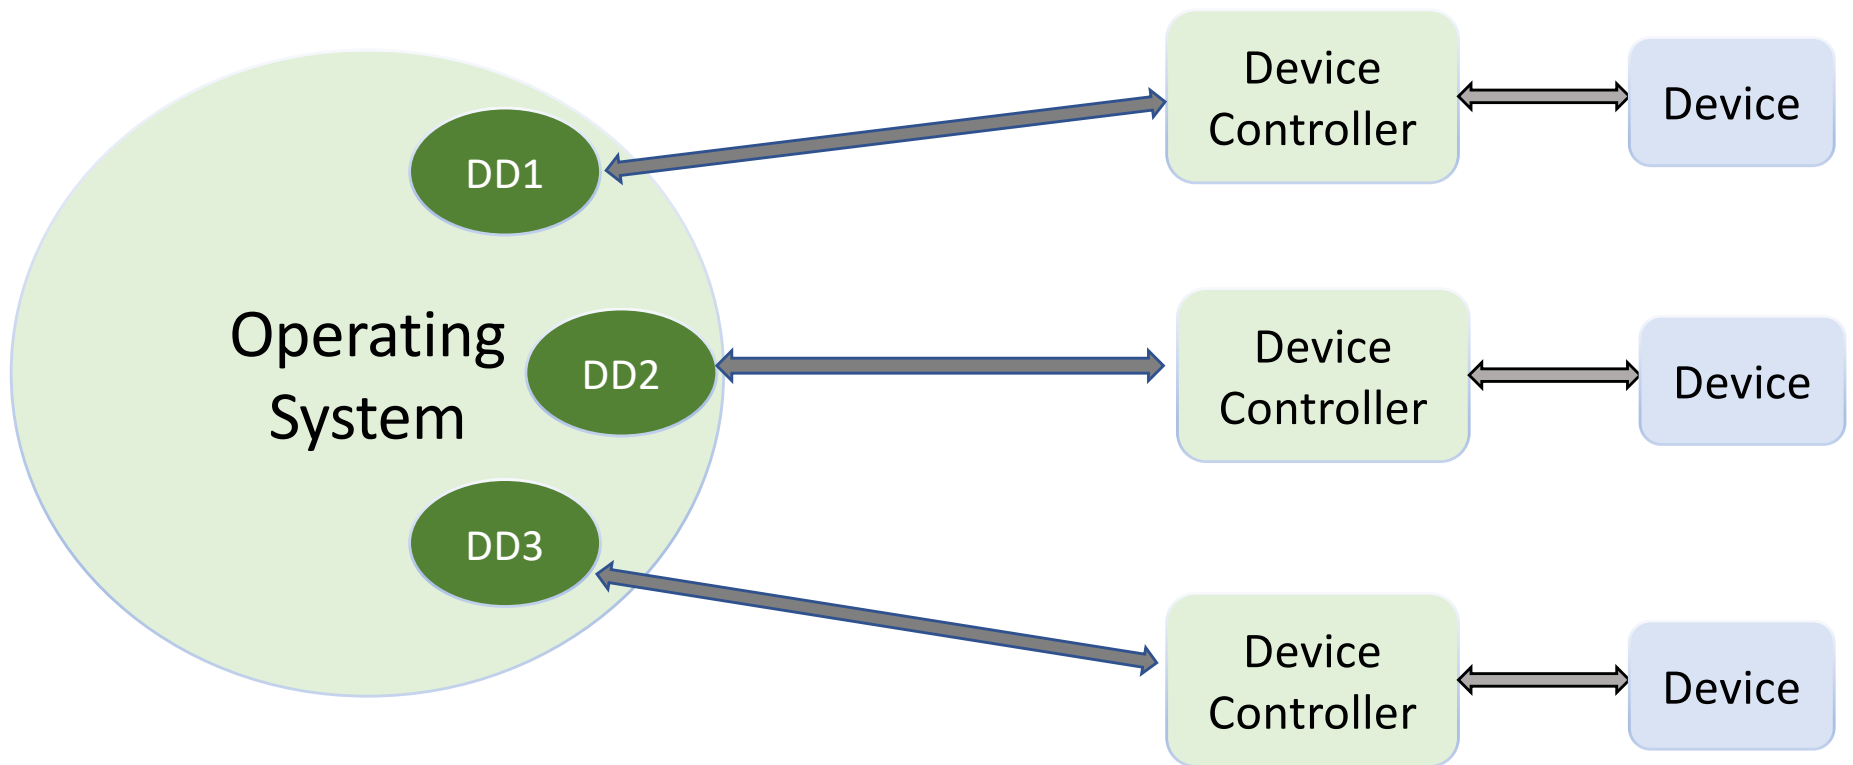
Computer System Operation: I/O Structure

- **Device driver:** An *operating system component* that provides uniform access to various devices and manages I/O to those devices.
- In an operating System there is a *device driver (DD)* for each device controller.



Computer System Operation: I/O Structure

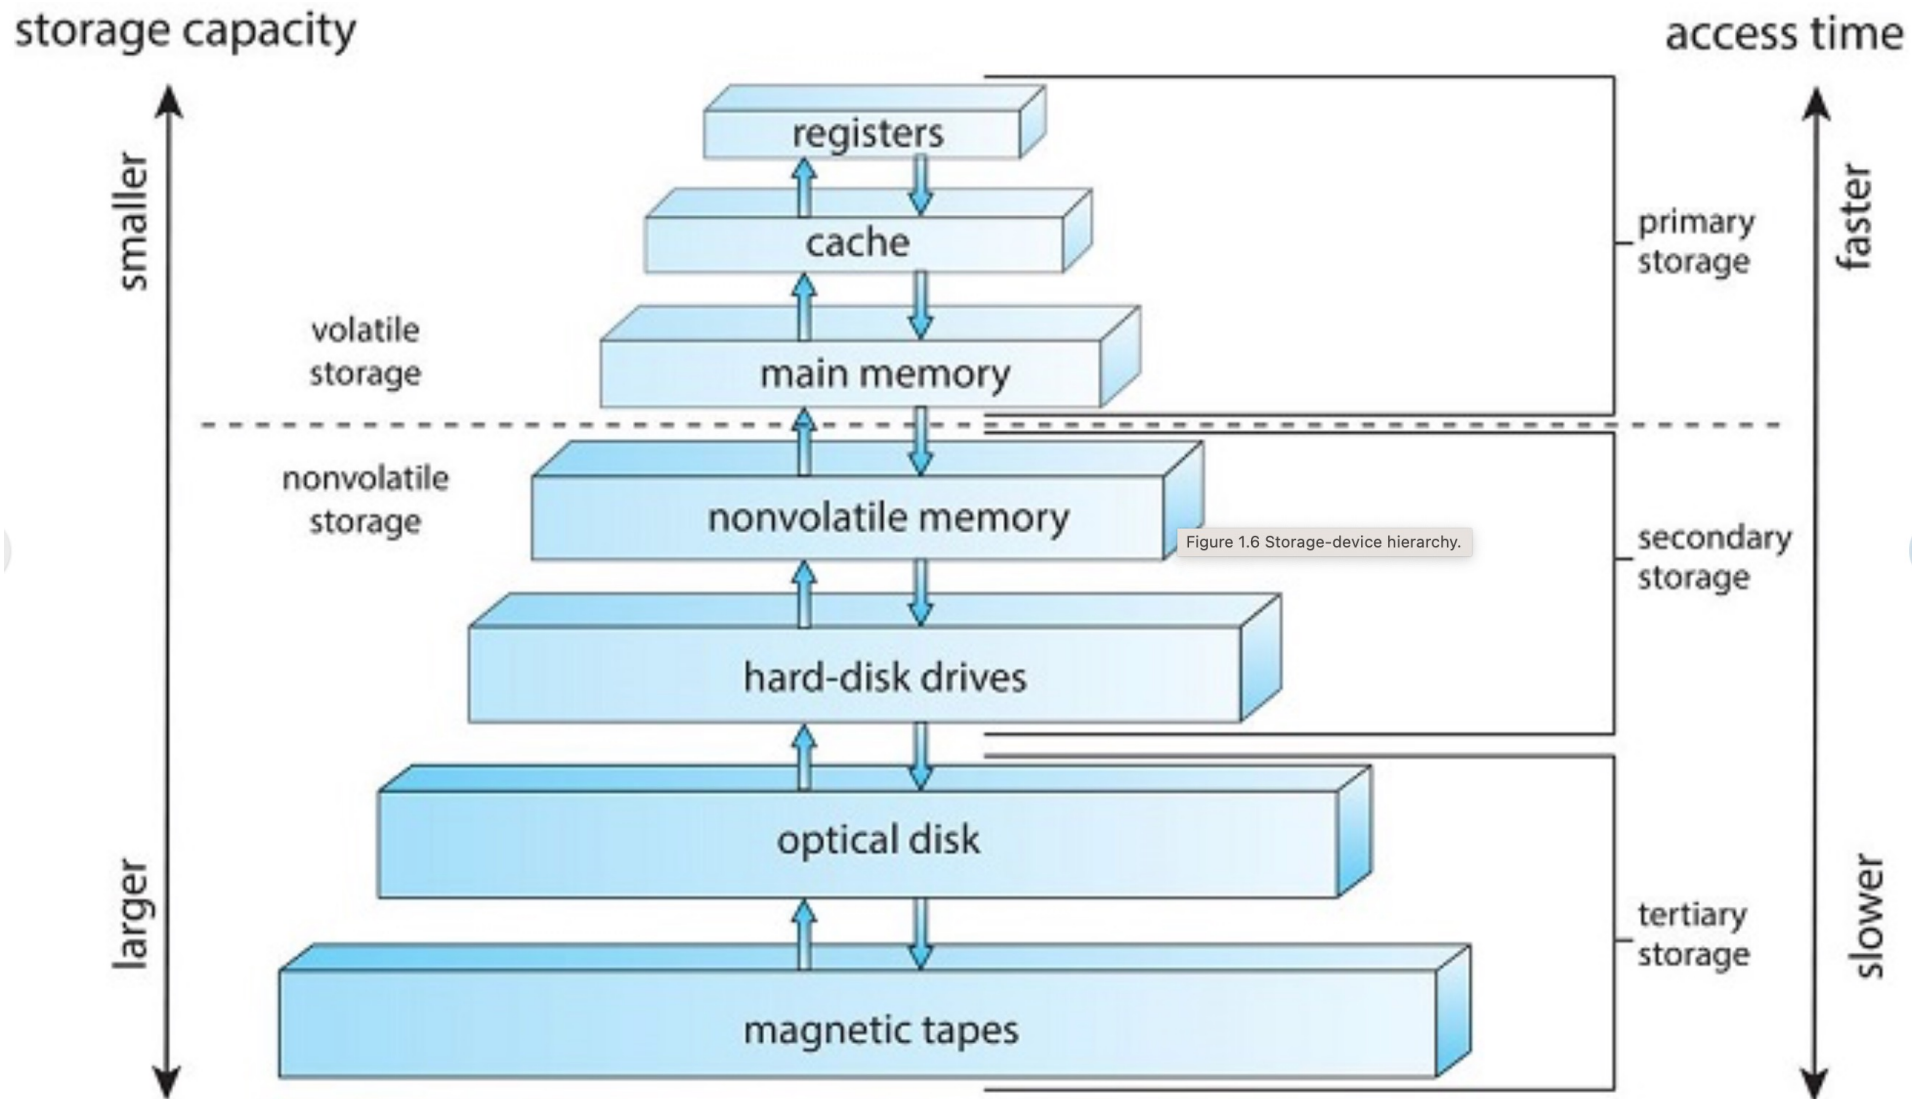
- Data movement between devices and the computer happens through the interaction between the device drivers and device controllers.



Computer System Operation

- CPU can access only data and instructions only from main memory (RAM)
- Programs must be brought into main memory for execution
 - Main memory is volatile
 - Secondary storage needed for permanent storage
- All forms of main memory provide an array of bytes, and each byte has its own address.

Computer System Operation: Storage-Device Hierarchy



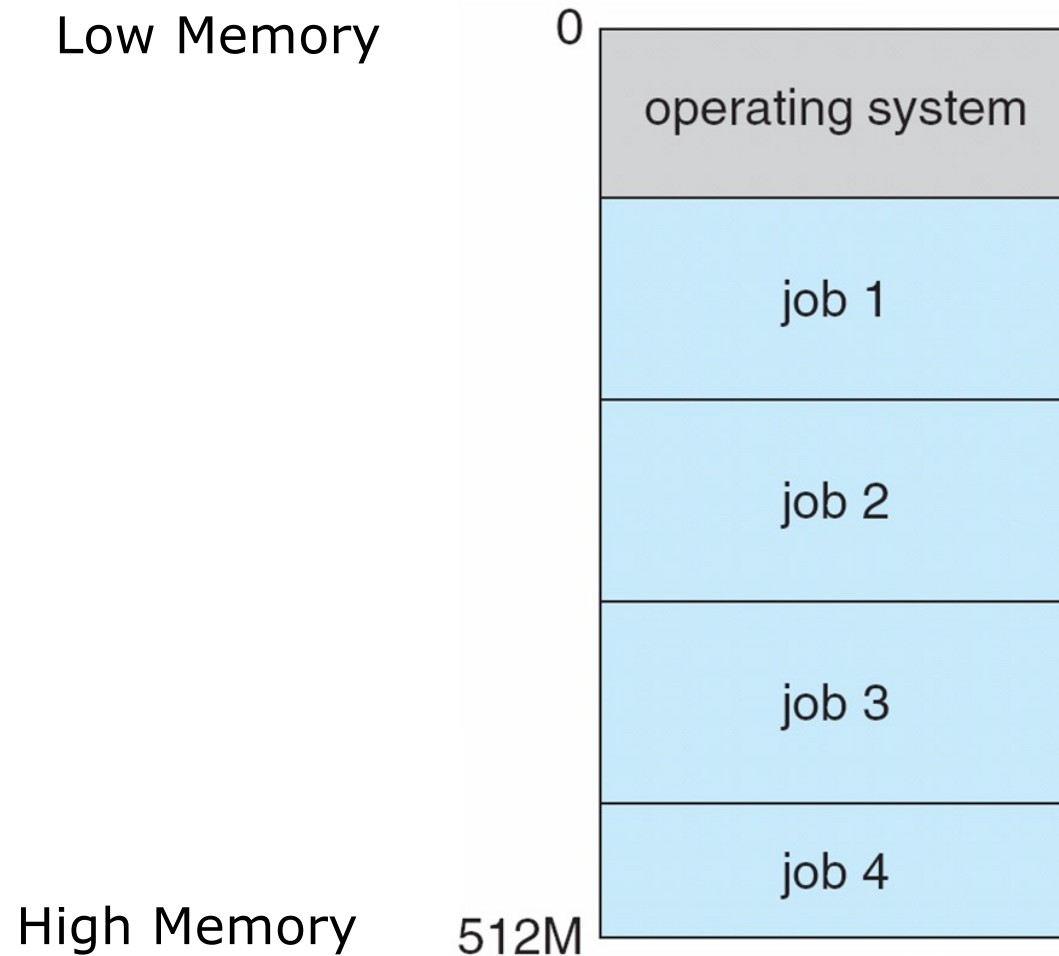
Operating System Operation: Multiprogramming

- An operating system provides the environment within which **multiple** programs are executed

How does multiprogramming work?

- OS organizes jobs so CPU always has one to execute
 - A subset of total jobs in system is kept in main memory
 - One job is selected and run via job scheduling
 - When it has to wait (for I/O for example), OS switches to another job

Memory Layout for Multiprogrammed System



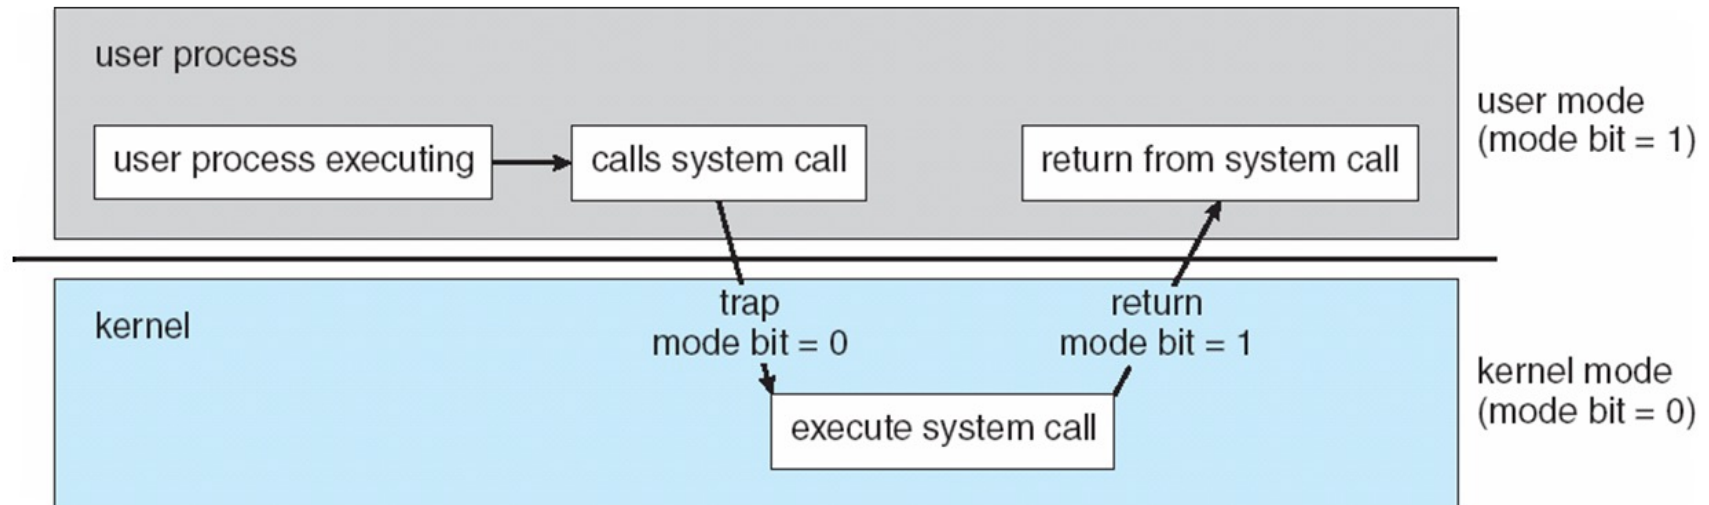
Operating System Operation

- **Timesharing** (**multitasking**) is logical extension of multiprogramming
 - CPU switches jobs so frequently that users can interact with each job while it is running, creating ***interactive*** computing
 - **Response time is < 1 second**

Operating-System Operations

- **Dual-mode** operation allows OS to protect itself and other system components
 - **User mode** and **kernel mode**
 - **Mode bit** is CPU Status bit used to indicate the current mode
 - Provides ability to distinguish when CPU is running in user code or kernel code
 - **Kernel mode bit = 0** and **User mode bit = 1**
 - **Privileged instructions** only executable in kernel mode
 - System call changes mode to kernel, return from call resets it to user (see next slide)

Transition from User to Kernel Mode - Illustration




Computer-System Architecture

- A **single processor** system has
 - One general purpose CPU
 - Possibly more than one special purpose CPUs (cannot process user request)
- **Multiprocessors** (parallel systems, multi-core systems)
 - More than one general purpose CPU's (sometimes share memory, bus and peripherals)
 - Advantages include: Increased throughput, Economy of scale, and Increased reliability

Multi-core Design

- Multiprocessor systems can be Multi-chip and/or **multicore** (More than one core on a single chip)
- Multicore systems are efficient as on chip communication is faster and consumes less power

Example of a dual core design:



Key functionality of an OS as a resource Manager

- Process Management
- Memory Management
- Storage and File management
- Protection and Security

Process

- Process is a program in execution
- Program is *passive* entity stored on disk (*executable file*), process is *active*
- Program becomes process when executable file loaded into memory

Process representation in Linux

- Processes in Linux are referred to as **tasks**.

- Represented by the C structure `task_struct`

```
pid_t pid; /* process identifier */
long state; /* state of the process */
unsigned int time_slice /* scheduling information */
struct task_struct *parent; /* this process's parent */
struct list_head children; /* this process's children */
struct files_struct *files; /* list of open files */
struct mm_struct *mm; /* address space of this process */
```

Kernel Data Structures

- Many similar to standard programming data structures
 - Linked lists (single, double and circular)
 - Binary search tree
 - Hash Tables
- Bitmap - string of n binary digits representing the status of n items
- Linux data structures defined in ***include*** files `<linux/list.h>`,
`<linux/kfifo.h>`, `<linux/rbtree.h>`

Computing Environments - Virtualization

- **Virtualization** is a *technology* that creates abstraction of the computer hardware.
- Thereby creating *an illusion* that all its operating systems are running on its own private computer.
- **VMM** (virtual machine Manager) provides virtualization services

Computing Environments - Virtualization

